

Supplementary Material for: Generalization and Computation for Policy Classes of Generative Adversarial Imitation Learning

A Analysis of Generalization Properties for Policy Classes

A.1 Proof of Theorem 1

Theorem 1. (Generalization for policy classes of distribution error)
Suppose Assumption 1 holds, and for the given D_I and π_E with

$$e(\hat{c}_{D_I}^{(m)} \rho^{\pi_E}, \hat{c}_{D_I}^{(m)} \rho^{\Pi}) \geq \sup_{D \in \mathcal{D}} \{e(\hat{c}_D^{(m)} \rho^{\pi_E}, \hat{c}_D^{(m)} \rho^{\Pi})\} - \hat{e},$$

then

$$\begin{aligned} & e(C_{D_I} \rho^{\pi_E}, C_{D_I} \rho^{\Pi}) \\ & \geq \underbrace{\sup_{D \in \mathcal{D}} \{e(\hat{c}_D^{(m)} \rho^{\pi_E}, \hat{c}_D^{(m)} \rho^{\Pi})\}}_{\text{Appr}(\mathcal{D}, m)} - \underbrace{2\hat{\mathfrak{A}}_{D_I}^{(m)}(C_{D_I} \rho^{\Pi}) - 8B_{\Pi} \sqrt{\frac{\log(3/\delta)}{2m}}}_{\text{Estm}(\Pi, m, \delta)} - \hat{e} \end{aligned}$$

for all $\delta \in (0, 1)$, with probability at least $1 - \delta$.

Proof. By (2) and Definition 2, we only need to prove

$$e(C_{D_I} \rho^{\pi_E}, C_{D_I} \rho^{\Pi}) - e(\hat{c}_{D_I}^{(m)} \rho^{\pi_E}, \hat{c}_{D_I}^{(m)} \rho^{\Pi})$$

has a lower bound.

$$\begin{aligned} & e(C_{D_I} \rho^{\pi_E}, C_{D_I} \rho^{\Pi}) - e(\hat{c}_{D_I}^{(m)} \rho^{\pi_E}, \hat{c}_{D_I}^{(m)} \rho^{\Pi}) \\ & = C_{D_I} \inf_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [\rho^{\pi_E}(s, a) - \rho^{\pi}(s, a)]\} \\ & \quad - \hat{c}_{D_I}^{(m)} \inf_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [\rho^{\pi_E}(s, a) - \rho^{\pi}(s, a)]\} \\ & \geq C_{D_I} \inf_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [\rho^{\pi_E}(s, a) - \rho^{\pi}(s, a)] - \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [\rho^{\pi_E}(s, a) - \rho^{\pi}(s, a)]\} \\ & \geq C_{D_I} \inf_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [\rho^{\pi_E}(s, a)] - \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [\rho^{\pi_E}(s, a)]\} \\ & \quad + C_{D_I} \inf_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [\rho^{\pi}(s, a)] - \mathbb{E}_{(s,a) \sim \mathbb{D}_I} [\rho^{\pi}(s, a)]\} \\ & = \left(\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^{\pi_E}(s, a)] - \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [C_{D_I} \rho^{\pi_E}(s, a)] \right) \\ & \quad - \sup_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^{\pi}(s, a)] - \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [C_{D_I} \rho^{\pi}(s, a)]\}. \end{aligned} \tag{13}$$

For notational simplicity, we denote $z_i = (s_{D_I}^{(i)}, a_{D_I}^{(i)})$, $Z = (z_1, \dots, z_m)$. Then

$$\begin{aligned}\mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [C_{D_I} \rho^{\pi_E}(s, a)] &= \hat{\mathbb{E}}_Z [C_{D_I} \rho^{\pi_E}], \\ \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [C_{D_I} \rho^\pi(s, a)] &= \hat{\mathbb{E}}_Z [C_{D_I} \rho^\pi].\end{aligned}$$

First, we show that $\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^{\pi_E}(s, a)] - \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [C_{D_I} \rho^{\pi_E}(s, a)]$ has a lower bound. Let

$$\phi_{\pi_E}(Z) = \mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^{\pi_E}(s, a)] - \hat{\mathbb{E}}_Z [C_{D_I} \rho^{\pi_E}].$$

Let Z and Z' be two samples differing by exactly one point, say $z_i \in Z, z'_i \in Z'$. Note that for all $(s, a) \in \mathcal{S} \times \mathcal{A}$, we have $C_{D_I} \rho^{\pi_E}(s, a) \leq B_\Pi$. Then

$$\begin{aligned}\phi_{\pi_E}(Z') - \phi_{\pi_E}(Z) &= (\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^{\pi_E}(s, a)] - \hat{\mathbb{E}}_{Z'} [C_{D_I} \rho^{\pi_E}]) \\ &\quad - (\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^{\pi_E}(s, a)] - \hat{\mathbb{E}}_Z [C_{D_I} \rho^{\pi_E}]) \\ &= \frac{1}{m} (C_{D_I} \rho^{\pi_E}(z_i) - C_{D_I} \rho^{\pi_E}(z'_i)) \leq \frac{2}{m} B_\Pi.\end{aligned}$$

By a similar derivation, we obtain $\phi_{\pi_E}(Z) - \phi_{\pi_E}(Z') \leq \frac{2}{m} B_\Pi$. Therefore, we have $|\phi_{\pi_E}(Z) - \phi_{\pi_E}(Z')| \leq \frac{2}{m} B_\Pi$. According to McDiarmid's inequality [26], we have

$$\phi_{\pi_E}(Z) - \mathbb{E}[\phi_{\pi_E}(Z)] \geq -2B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}}$$

with probability at least $1 - \frac{\delta}{3}$, where the outer expectation is taken over the random choice of $\hat{\mathbb{D}}_I$ with m state-action pairs. By the fact that $\mathbb{E}[\phi_{\pi_E}(Z)] = 0$, we have

$$\phi_{\pi_E}(Z) \geq -2B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} \quad (14)$$

with probability at least $1 - \frac{\delta}{3}$.

Next, we show that $\sup_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^\pi(s, a)] - \mathbb{E}_{(s,a) \sim \hat{\mathbb{D}}_I} [C_{D_I} \rho^\pi(s, a)]\}$ has an upper bound. Let

$$\phi_\pi(Z) = \sup_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_I} \rho^\pi(s, a)] - \hat{\mathbb{E}}_Z [C_{D_I} \rho^\pi]\}.$$

Note that for all $\pi \in \Pi$, we have $\max_{(s,a) \in \mathcal{S} \times \mathcal{A}} \{C_{D_1} \rho^\pi(s, a)\} \leq B_\Pi$. Then

$$\begin{aligned}
& \phi_\pi(Z') - \phi_\pi(Z) \\
&= (\sup_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_1} \rho^\pi(s, a)] - \hat{\mathbb{E}}_{Z'} [C_{D_1} \rho^\pi]\}) \\
&\quad - (\sup_{\pi \in \Pi} \{\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_1} \rho^\pi(s, a)] - \hat{\mathbb{E}}_Z [C_{D_1} \rho^\pi]\}) \\
&\leq \sup_{\pi \in \Pi} \left\{ (\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_1} \rho^\pi(s, a)] - \hat{\mathbb{E}}_{Z'} [C_{D_1} \rho^\pi]) \right. \\
&\quad \left. - (\mathbb{E}_{(s,a) \sim \mathbb{D}_I} [C_{D_1} \rho^\pi(s, a)] - \hat{\mathbb{E}}_Z [C_{D_1} \rho^\pi]) \right\} \\
&= \sup_{\pi \in \Pi} \{\hat{\mathbb{E}}_Z [C_{D_1} \rho^\pi] - \hat{\mathbb{E}}_{Z'} [C_{D_1} \rho^\pi]\} \\
&= \sup_{\pi \in \Pi} \left\{ \frac{C_{D_1} \rho^\pi(z_i) - C_{D_1} \rho^\pi(z'_i)}{m} \right\} \leq \frac{2}{m} B_\Pi.
\end{aligned}$$

By a similar derivation, we obtain that $\phi_\pi(Z) - \phi_\pi(Z') \leq \frac{2}{m} B_\Pi$. Therefore, $|\phi_\pi(Z) - \phi_\pi(Z')| \leq \frac{2}{m} B_\Pi$. According to McDiarmid's inequality,

$$\phi_\pi(Z) \leq \mathbb{E}[\phi_\pi(Z)] + 2B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} \quad (15)$$

with probability at least $1 - \frac{\delta}{3}$, where the outer expectation is taken over the random choice of $\hat{\mathbb{D}}_I$ with m state-action pairs.

In what follows, we prove that the right-hand side of (15) has an upper bound. By the standard Rademacher complexity technique [26], we obtain

$$\begin{aligned}
\mathbb{E}[\phi_\pi(Z)] &= \mathbb{E} \left[\sup_{\pi \in \Pi} \{\mathbb{E}_{Z''} [\hat{\mathbb{E}}_{Z''} [C_{D_1} \rho^\pi]] - \hat{\mathbb{E}}_Z [C_{D_1} \rho^\pi]\} \right] \\
&= \mathbb{E} \left[\sup_{\pi \in \Pi} \{\mathbb{E}_{Z''} [\hat{\mathbb{E}}_{Z''} [C_{D_1} \rho^\pi] - \hat{\mathbb{E}}_Z [C_{D_1} \rho^\pi]] \} \right] \\
&\leq \mathbb{E}_{Z, Z''} \left[\sup_{\pi \in \Pi} \{\hat{\mathbb{E}}_{Z''} [C_{D_1} \rho^\pi] - \hat{\mathbb{E}}_Z [C_{D_1} \rho^\pi]\} \right] \\
&= \mathbb{E}_{Z, Z''} \left[\frac{1}{m} \sup_{\pi \in \Pi} \left\{ \sum_{i=1}^m (C_{D_1} \rho^\pi(z''_i) - C_{D_1} \rho^\pi(z_i)) \right\} \right] \\
&= \mathbb{E}_{\sigma, Z, Z''} \left[\frac{1}{m} \sup_{\pi \in \Pi} \left\{ \sum_{i=1}^m \sigma_i (C_{D_1} \rho^\pi(z''_i) - C_{D_1} \rho^\pi(z_i)) \right\} \right] \\
&\leq \mathbb{E}_{\sigma, Z''} \left[\frac{1}{m} \sup_{\pi \in \Pi} \left\{ \sum_{i=1}^m \sigma_i C_{D_1} \rho^\pi(z''_i) \right\} \right] \\
&\quad + \mathbb{E}_{\sigma, Z} \left[\frac{1}{m} \sup_{\pi \in \Pi} \left\{ - \sum_{i=1}^m \sigma_i C_{D_1} \rho^\pi(z_i) \right\} \right] \\
&= 2\mathfrak{R}_{\mathbb{D}_I}^{(m)}(C_{D_1} \rho^\Pi), \quad (16)
\end{aligned}$$

where σ_i 's are i.i.d. Rademacher random variables.

According to McDiarmid's inequality,

$$\mathfrak{R}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi) \leq \hat{\mathfrak{R}}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi) + 2B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} \quad (17)$$

with probability at least $1 - \frac{\delta}{3}$, where

$$\hat{\mathfrak{R}}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi) = \mathbb{E}_\sigma \left[\frac{1}{m} \sup_{\pi \in \Pi} \left\{ \sum_{i=1}^m \sigma_i C_{D_I} \rho^\pi(z_i) \right\} \right].$$

Combining (15) with (16) and (17),

$$\phi_\pi(Z) \leq 2\hat{\mathfrak{R}}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi) + 6B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} \quad (18)$$

with probability at least $1 - \frac{2\delta}{3}$. Combining (14) with (18), we have

$$\phi_{\pi_E}(Z) - \phi_\pi(Z) \geq -2\hat{\mathfrak{R}}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi) - 8B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} \quad (19)$$

with probability at least $1 - \delta$. Combining (13) with (19) and the condition of Theorem 1, we complete the proof. \square

A.2 The lower bound of $e(C_{D_I}\rho^{\pi_E}, C_{D_I}\rho^\Pi)$ in terms of the covering number

The covering number of the function class $C_{D_I}\rho^\Pi$ under the ℓ_∞ distance $\|\cdot\|_\infty$ can be denoted as $\mathcal{N}(C_{D_I}\rho^\Pi, \epsilon, \|\cdot\|_\infty)$.

Corollary 2. *Under the same assumption of Theorem 1, then*

$$\begin{aligned} & e(C_{D_I}\rho^{\pi_E}, C_{D_I}\rho^\Pi) \\ & \geq \sup_{D \in \mathcal{D}} \{e(\hat{c}_D^{(m)}\rho^{\pi_E}, \hat{c}_D^{(m)}\rho^\Pi)\} - \frac{8}{m} - \frac{24B_\Pi}{\sqrt{m}} \sqrt{\log(\mathcal{N}(C_{D_I}\rho^\Pi, \frac{1}{\sqrt{m}}, \|\cdot\|_\infty))} \\ & \quad - 8B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} - \epsilon \end{aligned}$$

with probability at least $1 - \delta$.

Proof. We apply Dudley's entropy integral to the bound $\mathfrak{R}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi)$. Specifically, we have

$$\begin{aligned} \hat{\mathfrak{R}}_{\mathbb{D}_I}^{(m)}(C_{D_I}\rho^\Pi) & \leq \frac{4}{m} + \frac{12}{m} \int_{\frac{1}{\sqrt{m}}}^{\sqrt{m}B_\Pi} \sqrt{\log(\mathcal{N}(C_{D_I}\rho^\Pi, \epsilon, \|\cdot\|_\infty))} d\epsilon \\ & = \frac{4}{m} + \frac{12}{m} \sqrt{\log(\mathcal{N}(C_{D_I}\rho^\Pi, \xi_\pi, \|\cdot\|_\infty))} (\sqrt{m}B_\Pi - \frac{1}{\sqrt{m}}) \\ & \leq \frac{4}{m} + \frac{12B_\Pi}{\sqrt{m}} \sqrt{\log(\mathcal{N}(C_{D_I}\rho^\Pi, \frac{1}{\sqrt{m}}, \|\cdot\|_\infty))}, \end{aligned}$$

where $\xi_\pi \in (\frac{1}{\sqrt{m}}, \sqrt{m} \frac{B_\Pi}{C_{D_1}})$. Plugging this into Theorem 1, we obtain

$$\begin{aligned} & e(C_{D_1}\rho^{\pi_E}, C_{D_1}\rho^\Pi) \\ & \geq \sup_{D \in \mathcal{D}} \{e(\hat{c}_D^{(m)} \rho^{\pi_E}, \hat{c}_D^{(m)} \rho^\Pi)\} - \frac{8}{m} - \frac{24B_\Pi}{\sqrt{m}} \sqrt{\log(\mathcal{N}(C_{D_1}\rho^\Pi, \frac{1}{\sqrt{m}}, \|\cdot\|_\infty))} \\ & \quad - 8B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} - \hat{\epsilon} \end{aligned}$$

with probability at least $1 - \delta$. \square

A.3 Proof of Corollary 1

Corollary 1. *Suppose Assumption 2 holds and $\|\theta\|_2 \leq B_\theta$. Then*

$$\begin{aligned} e(C_{D_1}\rho^{\pi_E}, C_{D_1}\rho^\Pi) & \geq \sup_{D \in \mathcal{D}} \{e(\hat{c}_D^{(m)} \rho^{\pi_E}, \hat{c}_D^{(m)} \rho^\Pi)\} \\ & \quad - \frac{8}{m} - \frac{24B_\Pi}{\sqrt{m}} \sqrt{p \log(1 + 2\sqrt{2m}B_\theta L_h)} - 8B_\Pi \sqrt{\frac{\log(3/\delta)}{2m}} - \hat{\epsilon} \end{aligned}$$

with probability at least $1 - \delta$.

Proof. $C_{D_1}\rho^\pi(s, a)$ can be bounded by

$$|C_{D_1}\rho^\pi(s, a)| = |\theta^\top h(\psi_s, \psi_a)| \leq \|\theta\|_2 \|h(\psi_s, \psi_a)\|_2 \leq \sqrt{2}B_\theta L_h,$$

where the first inequality comes from Cauchy-Schwartz inequality.

To compute the covering number, we exploit the Lipschitz continuity of $C_{D_1}\rho^\pi(s, a)$ with respect to the parameter θ . Specifically, for two different parameters θ and θ' , we have

$$\begin{aligned} & \|C_{D_1}\rho^{\tilde{\pi}_\theta}(\psi_s, \psi_a) - C_{D_1}\rho^{\tilde{\pi}_{\theta'}}(\psi_s, \psi_a)\|_\infty \\ & = \|(\theta - \theta')^\top h(\psi_s, \psi_a)\|_\infty \\ & \stackrel{(i)}{\leq} \|\theta - \theta'\|_2 \sup_{(s,a) \in \mathcal{S} \times \mathcal{A}} \{\|h(\psi_s, \psi_a)\|_2\} \\ & \stackrel{(ii)}{\leq} \|\theta - \theta'\|_2 L_h \sup_{(s,a) \in \mathcal{S} \times \mathcal{A}} \{\sqrt{\|\psi_s\|_2^2 + \|\psi_a\|_2^2}\} \\ & \stackrel{(iii)}{\leq} \sqrt{2}L_h \|\theta - \theta'\|_2, \end{aligned}$$

where (i) comes from Cauchy-Schwartz inequality, (ii) comes from the Lipschitz continuity of h , and (iii) comes from the boundedness of ψ_s and ψ_a .

Denote $\Theta = \{\theta : \|\theta\|_2 \leq B_\theta\}$. Then we have

$$\begin{aligned} \mathcal{N}(C_{D_I}\rho^\Pi, \frac{1}{\sqrt{m}}, \|\cdot\|_\infty) &\leq \mathcal{N}(\Theta, \frac{1}{\sqrt{2m}L_h}, \|\cdot\|_2) \\ &\leq \left(1 + \frac{2B_\theta}{\frac{1}{\sqrt{2m}L_h}}\right)^p \\ &= (1 + 2\sqrt{2m}B_\theta L_h)^p. \end{aligned} \quad (20)$$

Plugging (20) into Theorem 1, we complete the proof. \square

A.4 Proof of Theorem 2

Theorem 2. (*GAIL Generalization for policy classes*) *Under the same assumption of Theorem 1, we have*

$$V_{\pi_E} - \sup_{\pi \in \Pi} V_\pi \geq \frac{1}{1-\gamma} (\text{Appr}(\mathcal{D}, m) + \text{Estm}(\Pi, m, \delta) - \hat{\epsilon})$$

with probability at least $1 - \delta$.

Proof. Combining the definition of V_π in (1) with the reward function of GAIL $r_{D_I}(s, a) = -\log(1 - D_I(s, a))$ [20,45], we obtain

$$\begin{aligned} &V_{\pi_E} - \sup_{\pi \in \Pi} V_\pi \\ &= \frac{1}{1-\gamma} \mathbb{E}_{(s,a) \sim \rho^{\pi_E}} [r(s, a)] - \sup_{\pi \in \Pi} \left\{ \frac{1}{1-\gamma} \mathbb{E}_{(s,a) \sim \rho^\pi} [r(s, a)] \right\} \\ &= \frac{1}{1-\gamma} \inf_{\pi \in \Pi} \left\{ \mathbb{E}_{(s,a) \sim \rho^{\pi_E}} [-\log(1 - D_I(s, a))] \right. \\ &\quad \left. - \mathbb{E}_{(s,a) \sim \rho^\pi} [-\log(1 - D_I(s, a))] \right\} \\ &= \frac{1}{1-\gamma} e(C_{D_I}\rho^{\pi_E}, C_{D_I}\rho^\Pi). \end{aligned} \quad (21)$$

Plugging (21) into the conclusion of Theorem 1, we complete the proof. \square

B Algorithm Diagrams

To better illustrate the computational algorithms for (3), we summarize the SAM submodule in Algorithm 1 and the TSSG algorithm in Algorithm 2. The framework of TSSG is shown in Figure 3.

C Convergence Analysis

The convergence of soft policy iteration in SAC algorithms has been proved in [19]. Therefore, for a fixed ϕ , the SAM submodule converges. Then we only need to prove the convergence of Algorithm 2.

First, the definition of stationary point and two assumptions are given:

Algorithm 1 SAM submodule: $\theta_T = \text{SAM}(\theta_0; w_0, \alpha_0, \phi)$

Input: $\theta_0 \in \mathbb{R}^p$, $w_0 \in \mathbb{R}^d$, $\alpha_0 \in \mathbb{R}$, fixed $\phi \in \mathbb{R}^q$;**Output:** θ_T .

- 1: **for** $t = 0, \dots, T - 1$ **do**
 - 2: Apply (9) and (10) to update w_{t+1} and θ_{t+1} respectively.
 - 3: Apply (12) to automate entropy adjustment for α_{t+1} .
 - 4: **end for**
-

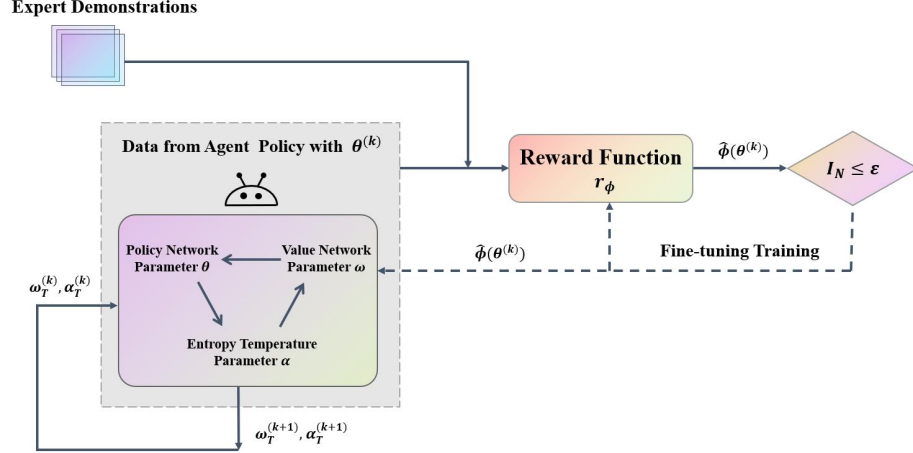
Algorithm 2 TSSG

Input: $\theta^{(0)} \in \mathbb{R}^p$, $w^{(0)} \in \mathbb{R}^d$, $\alpha_0 \in \mathbb{R}$;**Output:** $\theta^{(N)}$, $\phi^{(N)}$.

- 1: **for** $k = 0, \dots, N - 1$ **do**
- 2: Apply (11) with samples collected by $\pi_{\theta^{(k)}}$ and π_E to update $\hat{\phi}(\theta^{(k)})$.
- 3: Apply the SAM submodule with samples collected by $\pi_{\theta^{(k)}}$ to update the policy:

$$\theta^{(k+1)} = \text{SAM}(\theta^{(k)}; w_T^{(k)}, \alpha_T^{(k)}, \hat{\phi}(\theta^{(k)})).$$

- 4: **end for**
-

**Fig. 3.** The TSSG framework for solving GAIL.

Definition 3. (*stationary point*) *There exists T such that the SAM submodule converges and $(\theta^*, w_0^*, \alpha_0^*)$ is a stationary point for the SAM submodule, then θ^* is a stationary point for TSSG if*

$$\sum_{t=0}^{T-1} \nabla_{\theta} F(\theta_t^*; w_{t+1}^*, \phi^*(\theta^*), \alpha_t^*) = 0,$$

where $\theta_0^* = \theta^*$.

The stationarity is a necessary condition for optimality. Accordingly, the formula below can be used to measure the sub-stationarity of the algorithm at the N -th iteration:

$$I_N = \min_{0 \leq k \leq N-1} \left(\mathbb{E} \left\| \sum_{t=0}^{T-1} \nabla_{\theta} F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right\|_2^2 \right).$$

Assumption 4. *For any $\theta, w, \alpha, \hat{\phi}(\theta)$ and $k, t \in \mathbb{N}$, there exist constants $M_{\theta}, M_w, M_{\alpha} > 0$, subject to:*

Unbiasedness:

$$\mathbb{E}[\nabla_{\theta} \tilde{f}^{(j)}(\theta; w, \hat{\phi}(\theta), \alpha)] = \nabla_{\theta} F(\theta; w, \hat{\phi}(\theta), \alpha),$$

Gradient boundedness:

$$\begin{aligned} \mathbb{E}[\|\nabla_{\theta} \tilde{f}^{(j)}(\theta; w, \hat{\phi}(\theta), \alpha)\|_2^2] &\leq M_{\theta}, \\ \mathbb{E}[\|\nabla_w \tilde{J}_Q^{(j)}(w; \theta, \hat{\phi}(\theta), \alpha)\|_2^2] &\leq M_w, \\ \mathbb{E}[\|\nabla_{\alpha} \tilde{J}_t^{(k)}(\alpha; \theta)\|_2^2] &\leq M_{\alpha}. \end{aligned}$$

Assumption 4 is mild and satisfied by first-order optimization algorithms [28,9,13,15], e.g., the greedy stochastic gradient algorithm [10].

Assumption 5. (1) *There exists a constant $L_Q > 0$, subject to for any w and w' , we have*

$$\|Q_w^{\text{soft}} - Q_{w'}^{\text{soft}}\|_{\infty} \leq L_Q \|w - w'\|_2.$$

(2) *There exists constants $S_g, S_Q > 0$, subject to for any θ and θ' , we have*

$$\begin{aligned} \|\mathbb{E}_{\rho^{\pi_{\theta}}} [g(\psi_s, \psi_a)] - \mathbb{E}_{\rho^{\pi_{\theta'}}} [g(\psi_s, \psi_a)]\|_2 &\leq S_g \|\theta - \theta'\|_2, \\ \|\nabla_{\theta} \mathbb{E}_{\rho^{\pi_{\theta}}} [Q_w^{\text{soft}}(\psi_s, \psi_a; \theta, \phi)] - \nabla_{\theta} \mathbb{E}_{\rho^{\pi_{\theta'}}} [Q_w^{\text{soft}}(\psi_s, \psi_a; \theta', \phi)]\|_2 &\leq S_Q \|\theta - \theta'\|_2. \end{aligned}$$

(3) *There exists constant $B_H, L_{\alpha}, S_H > 0$, subject to for any α, α', θ and θ' , we have*

$$\begin{aligned} \mathbb{E}_{s \sim d^{\pi_{\theta}}} [\alpha \mathbb{H}(\pi_{\theta}(\cdot|s))] &\leq B_H, \\ |\mathbb{E}_{s \sim d^{\pi_{\theta}}} [\alpha \mathbb{H}(\pi_{\theta}(\cdot|s))] - \mathbb{E}_{s \sim d^{\pi_{\theta}}} [\alpha' \mathbb{H}(\pi_{\theta}(\cdot|s))]| &\leq L_{\alpha} \|\alpha - \alpha'\|_2, \\ \|\nabla_{\theta} \mathbb{E}_{s \sim d^{\pi_{\theta}}} [\alpha \mathbb{H}(\pi_{\theta}(\cdot|s))] - \nabla_{\theta} \mathbb{E}_{s \sim d^{\pi_{\theta'}}} [\alpha \mathbb{H}(\pi_{\theta}(\cdot|s))]\|_2 &\leq S_H \|\theta - \theta'\|_2. \end{aligned}$$

(1) of Assumption 5 is the Lipschitz continuity of the soft Q-function with respect to its parameter w . (2) characterizes some Lipschitz continuity conditions with respect to the parameter θ . (3) states some common regularity conditions for entropy [17,18,19,31]. The convergence of the TSSG algorithm is organized as follows.

Theorem 3. *Suppose Assumptions 3,4,5 hold. Given T , in the condition that the SAM submodule converges, for any $\epsilon > 0$, we take*

$$\begin{aligned}\eta_\theta &= \frac{\epsilon}{2T^2 M_\theta (2S_\theta + TS_g^2/\mu)}, \\ \eta_w &= \frac{\epsilon^2}{8T^4 L_Q M_\theta \sqrt{M_w} (2S_\theta + TS_g^2/\mu)}, \\ \eta_\alpha &= \frac{\epsilon^2}{8T^4 L_\alpha M_\theta \sqrt{M_\alpha} (2S_\theta + TS_g^2/\mu)},\end{aligned}$$

where $S_\theta = S_H + S_Q$. Then at most

$$N = \tilde{O}\left(\frac{T^3 B_F M_\theta (S_\theta + TS_g^2/\mu)}{\epsilon^2}\right)$$

iterations such that $I_N \leq \epsilon$, where $B_F = \frac{\sqrt{2}\kappa L_g(2-\gamma)}{1-\gamma} + \frac{B_H}{1-\gamma} + \frac{\mu}{2}\kappa^2$.

Here \tilde{O} hides high dependence on T and linear or quadratic dependence on some constants in Assumptions 3-5. Next, we prove Theorem 3.

C.1 Boundedness of soft Q-function

Lemma 1. *For any w , we have $\|Q_w^{\text{soft}}\|_\infty \leq B_Q$, where $B_Q = \frac{\sqrt{2}\kappa L_g}{1-\gamma} + \frac{\gamma B_H}{1-\gamma}$.*

Proof. The reward function $r_\phi(s, a)$ can be bounded by

$$|r_\phi(s, a)| \leq \|\phi\|_2 \cdot \|g(\psi_s, \psi_a)\|_2 \leq \sqrt{2}\kappa L_g.$$

By the definition of soft Q-function [17], we have

$$\begin{aligned}Q_w^{\text{soft}}(s_t, a_t; \theta, \phi) &= r_\phi(s_t, a_t) + \mathbb{E}_{(s_{t+1}, \dots) \sim d^{\pi_\theta}} \left[\sum_{l=1}^{\infty} \gamma^l (r_\phi(s_{t+l}, a_{t+l}) + \alpha \mathbb{H}(\pi_\theta(\cdot | s_{t+l}))) \right] \\ &= \sum_{l=0}^{\infty} \gamma^l r_\phi(s_{t+l}, a_{t+l}) + \sum_{l=1}^{\infty} \gamma^l \mathbb{E}_{s \sim d^{\pi_\theta}} [\alpha \mathbb{H}(\pi_\theta(\cdot | s_{t+l}))] \\ &\leq \sum_{l=0}^{\infty} \gamma^l \sqrt{2}\kappa L_g + \sum_{l=1}^{\infty} \gamma^l B_H \\ &\leq \frac{\sqrt{2}\kappa L_g}{1-\gamma} + \frac{\gamma B_H}{1-\gamma}.\end{aligned}$$

□

C.2 Lipschitz properties of the gradients

Lemma 2. *Suppose Assumption 5 holds. For any $\theta, \theta', w, \phi, \alpha$, we have*

$$\|\nabla_\theta F(\theta; w, \phi, \alpha) - \nabla_\theta F(\theta'; w, \phi, \alpha)\|_2 \leq S_\theta \|\theta - \theta'\|_2,$$

where $S_\theta = S_H + S_Q$.

Proof. By [18,19], we have

$$\begin{aligned} & \nabla_\theta F(\theta; w, \phi, \alpha) \\ &= \mathbb{E}_{s_t \sim d^{\pi_\theta}} [\alpha \nabla_\theta \log(\pi_\theta(a_t|s_t)) \\ & \quad + (\alpha \nabla_{a_t} \log(\pi_\theta(a_t|s_t)) - \nabla_{a_t} Q_w^{\text{soft}}(s_t, a_t; \theta, \phi)) \nabla_\theta a_t] \\ &= -\nabla_\theta \mathbb{E}_{s_t \sim d^{\pi_\theta}} [\alpha \mathbb{H}(\pi_\theta(\cdot|s_t))] - \nabla_\theta \mathbb{E}_{(s_t, a_t) \sim \rho^{\pi_\theta}} [Q_w^{\text{soft}}(s_t, a_t; \theta, \phi)]. \end{aligned}$$

Therefore,

$$\begin{aligned} & \|\nabla_\theta F(\theta; w, \phi, \alpha) - \nabla_\theta F(\theta'; w, \phi, \alpha)\|_2 \\ & \leq \|\nabla_\theta \mathbb{E}_{s_t \sim d^{\pi_\theta}} [\alpha \mathbb{H}(\pi_\theta(\cdot|s_t))] - \nabla_\theta \mathbb{E}_{s_t \sim d^{\pi_{\theta'}}} [\alpha \mathbb{H}(\pi_{\theta'}(\cdot|s_t))]\|_2 \\ & \quad + \|\nabla_\theta \mathbb{E}_{(s_t, a_t) \sim \rho^{\pi_\theta}} [Q_w^{\text{soft}}(s_t, a_t; \theta, \phi)] - \nabla_\theta \mathbb{E}_{(s_t, a_t) \sim \rho^{\pi_{\theta'}}} [Q_w^{\text{soft}}(s_t, a_t; \theta', \phi)]\|_2 \\ & \leq (S_H + S_Q) \|\theta - \theta'\|_2. \end{aligned}$$

□

C.3 Boundedness of F

Lemma 3. *Under Assumption 5, there exists $B_F = \frac{\sqrt{2}\kappa L_g(2-\gamma)}{1-\gamma} + \frac{B_H}{1-\gamma} + \frac{\mu}{2}\kappa^2$ such that for any θ, w, ϕ, α , we have $|F(\theta; w, \phi, \alpha)| \leq B_F$.*

Proof. By the definition of $F(\theta; w, \phi, \alpha)$, we have

$$\begin{aligned} & |F(\theta; w, \phi, \alpha)| \\ & \leq \mathbb{E}_{(s, a) \sim \rho^{\pi_\theta}} [r_\phi(s, a)] + \mathbb{E}_{s_t \sim d^{\pi_\theta}} [\alpha \mathbb{H}(\pi_\theta(\cdot|s_t))] + \|Q_w^{\text{soft}}\|_\infty + \frac{\mu}{2} \|\phi\|_2^2 \\ & \leq \sqrt{2}\kappa L_g + B_H + B_Q + \frac{\mu}{2} \kappa^2 \\ & = \frac{\sqrt{2}\kappa L_g(2-\gamma)}{1-\gamma} + \frac{B_H}{1-\gamma} + \frac{\mu}{2} \kappa^2. \end{aligned}$$

□

C.4 Proof of Theorem 3

Theorem 3. Suppose Assumptions 3,4,5 hold. Given T , in the condition that the SAM submodule converges, for any $\epsilon > 0$, we take

$$\begin{aligned}\eta_\theta &= \frac{\epsilon}{2T^2 M_\theta (2S_\theta + TS_g^2/\mu)}, \\ \eta_w &= \frac{\epsilon^2}{8T^4 L_Q M_\theta \sqrt{M_w} (2S_\theta + TS_g^2/\mu)}, \\ \eta_\alpha &= \frac{\epsilon^2}{8T^4 L_\alpha M_\theta \sqrt{M_\alpha} (2S_\theta + TS_g^2/\mu)},\end{aligned}$$

where $S_\theta = S_H + S_Q$. Then at most

$$N = \tilde{O}\left(\frac{T^3 B_F M_\theta (S_\theta + TS_g^2/\mu)}{\epsilon^2}\right)$$

iterations such that $I_N \leq \epsilon$, where $B_F = \frac{\sqrt{2}\kappa L_g(2-\gamma)}{1-\gamma} + \frac{B_H}{1-\gamma} + \frac{\mu}{2}\kappa^2$.

Proof. Employing Lemma 2 and the Mean Value Theorem, we have

$$\begin{aligned}& \sum_{t=0}^{T-1} \left(F(\theta_{t+1}^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) - F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right) \\& - \sum_{t=0}^{T-1} \langle \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), \theta_{t+1}^{(k)} - \theta_t^{(k)} \rangle \\& = \sum_{t=0}^{T-1} \left(\langle \nabla_\theta F(\tilde{\theta}_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), \theta_{t+1}^{(k)} - \theta_t^{(k)} \rangle \right. \\& \quad \left. - \langle \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), \theta_{t+1}^{(k)} - \theta_t^{(k)} \rangle \right) \\& \leq \sum_{t=0}^{T-1} \left\| \nabla_\theta F(\tilde{\theta}_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right. \\& \quad \left. - \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right\|_2 \|\theta_{t+1}^{(k)} - \theta_t^{(k)}\|_2 \\& \leq S_\theta \sum_{t=0}^{T-1} \|\tilde{\theta}_t^{(k)} - \theta_t^{(k)}\|_2 \|\theta_{t+1}^{(k)} - \theta_t^{(k)}\|_2 \leq S_\theta \sum_{t=0}^{T-1} \|\theta_{t+1}^{(k)} - \theta_t^{(k)}\|_2^2, \tag{22}\end{aligned}$$

where $\tilde{\theta}_t^{(k)}$ is some interpolation between $\theta_{t+1}^{(k)}$ and $\theta_t^{(k)}$. Note that

$$\begin{aligned}& \mathbb{E} \langle \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), \theta_{t+1}^{(k)} - \theta_t^{(k)} \rangle \\& \stackrel{(i)}{=} \mathbb{E} \langle \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), -\eta_\theta (\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)}) + \xi_{\theta_t}^k) \rangle \\& \stackrel{(ii)}{=} \mathbb{E} \langle \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), -\eta_\theta \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)}) \rangle \\& \stackrel{(iii)}{=} \mathbb{E} \langle \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}), -\eta_\theta \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \rangle \\& \stackrel{(iv)}{=} -\eta_\theta \mathbb{E} \|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2, \tag{23}\end{aligned}$$

where

$$\xi_{\theta_t}^k = \frac{1}{n_\theta} \sum_{j \in D_\theta^t} \nabla_\theta \tilde{f}^{(j)}(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)}) - \nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)}).$$

Here (ii) comes from the unbiased property of $\tilde{f}^{(j)}$, and (iii) comes from the unbiased property of $\hat{\phi}(\theta^{(k)})$ and the linearity of $\nabla_\theta F$ in ϕ . Now taking the expectation on both sides of (22) and plugging (23) in, we obtain

$$\begin{aligned} & \sum_{t=0}^{T-1} \left(\mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) - \mathbb{E} F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right) \\ & + \eta_\theta \sum_{t=0}^{T-1} (\mathbb{E} \|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \\ & \leq S_\theta \left(\sum_{t=0}^{T-1} \mathbb{E} \|\theta_{t+1}^{(k)} - \theta_t^{(k)}\|_2^2 \right) \\ & = S_\theta \eta_\theta^2 \left(\sum_{t=0}^{T-1} \frac{1}{n_\theta^2} \mathbb{E} \left\| \sum_{j \in D_\theta^t} (\nabla_\theta \tilde{f}^{(j)}(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)})) \right\|_2^2 \right) \\ & \leq S_\theta \eta_\theta^2 \left(\sum_{t=0}^{T-1} \left(\frac{1}{n_\theta} \sum_{j \in D_\theta^t} \mathbb{E} \|\nabla_\theta \tilde{f}^{(j)}(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)})\|_2^2 \right) \right) \\ & \leq S_\theta \eta_\theta^2 T M_\theta. \end{aligned} \tag{24}$$

Dividing both sides by η_θ and rearranging the terms in (24), we get

$$\begin{aligned} & \sum_{t=0}^{T-1} (\mathbb{E} \|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \\ & \leq \frac{\sum_{t=0}^{T-1} \left(\mathbb{E} F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) - \mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+2}^{(k)}, \phi^*(\theta^{(k)}), \alpha_{t+1}^{(k)}) \right)}{\eta_\theta} \\ & + \frac{\sum_{t=0}^{T-1} \left(\mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+2}^{(k)}, \phi^*(\theta^{(k)}), \alpha_{t+1}^{(k)}) - \mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right)}{\eta_\theta} \\ & + S_\theta \eta_\theta T M_\theta \\ & = \frac{\mathbb{E} F(\theta^{(k)}; w_1^{(k)}, \phi^*(\theta^{(k)}), \alpha_0^{(k)}) - \mathbb{E} F(\theta^{(k+1)}; w_{T+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)})}{\eta_\theta} \\ & + \frac{\sum_{t=0}^{T-1} \left(\mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+2}^{(k)}, \phi^*(\theta^{(k)}), \alpha_{t+1}^{(k)}) - \mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right)}{\eta_\theta} \\ & + S_\theta \eta_\theta T M_\theta. \end{aligned} \tag{25}$$

By Assumption 5, for any $\theta, w, w', \theta', \alpha, \alpha'$, we have

$$\begin{aligned}
& |F(\theta; w, \phi^*(\theta'), \alpha) - F(\theta; w', \phi^*(\theta'), \alpha')| \\
&= |\mathbb{E}_{s_t \sim \mathcal{D}_1^{(k)}} [\mathbb{E}_{a_t \sim \pi_\theta} [\alpha \log(\pi_\theta(a_t|s_t)) - Q_w^{\text{soft}}(s_t, a_t; \theta, \phi^*(\theta'))]] \\
&\quad - \mathbb{E}_{s_t \sim \mathcal{D}_1^{(k)}} [\mathbb{E}_{a_t \sim \pi_\theta} [\alpha' \log(\pi_\theta(a_t|s_t)) - Q_{w'}^{\text{soft}}(s_t, a_t; \theta, \phi^*(\theta'))]]| \\
&\leq |\mathbb{E}_{s_t \sim \mathcal{D}_1^{(k)}} [\alpha \mathbb{H}(\pi_\theta(\cdot|s_t))] - \mathbb{E}_{s_t \sim \mathcal{D}_1^{(k)}} [\alpha' \mathbb{H}(\pi_\theta(\cdot|s_t))]| + \|Q_w^{\text{soft}} - Q_{w'}^{\text{soft}}\|_\infty \\
&\leq L_\alpha \|\alpha - \alpha'\| + L_Q \|w - w'\|_2.
\end{aligned}$$

Therefore, we obtain

$$\begin{aligned}
& \sum_{t=0}^{T-1} \left(\mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+2}^{(k)}, \phi^*(\theta^{(k)}), \alpha_{t+1}^{(k)}) - \mathbb{E} F(\theta_{t+1}^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)}) \right) \\
&\leq \frac{\sum_{t=0}^{T-1} \left(L_\alpha \mathbb{E} \|\alpha_{t+1}^{(k)} - \alpha_t^{(k)}\|_2 + L_Q \mathbb{E} \|w_{t+2}^{(k)} - w_{t+1}^{(k)}\|_2 \right)}{\eta_\theta} \\
&= \sum_{t=0}^{T-1} \left(L_\alpha \eta_\alpha \mathbb{E} \|\nabla_\alpha \tilde{J}_t^{(k)}(\alpha_t^{(k)}; \theta_{t+1}^{(k)})\|_2 \right. \\
&\quad \left. + L_Q \eta_w \frac{\sum_{j \in D_w^t} \mathbb{E} \|\nabla_w \tilde{J}_Q^j(w_{t+1}^{(k)}; \theta_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_{t+1}^{(k)})\|_2}{n_w} \right) \\
&\leq \sum_{t=0}^{T-1} \left(L_\alpha \eta_\alpha \sqrt{\mathbb{E} \|\nabla_\alpha \tilde{J}_t^{(k)}(\alpha_t^{(k)}; \theta_{t+1}^{(k)})\|_2^2} \right. \\
&\quad \left. + L_Q \eta_w \frac{\sum_{j \in D_w^t} \sqrt{\mathbb{E} \|\nabla_w \tilde{J}_Q^j(w_{t+1}^{(k)}; \theta_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_{t+1}^{(k)})\|_2^2}}{n_w} \right) \\
&\leq T L_\alpha \sqrt{M_\alpha} \eta_\alpha + T L_Q \sqrt{M_w} \eta_w. \tag{26}
\end{aligned}$$

Plugging (26) into (25), we get

$$\begin{aligned}
& \sum_{t=0}^{T-1} \left(\mathbb{E} \|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2 \right) \\
&\leq \frac{\mathbb{E} F(\theta^{(k)}; w_1^{(k)}, \phi^*(\theta^{(k)}), \alpha_0^{(k)}) - \mathbb{E} F(\theta^{(k+1)}; w_{T+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)})}{\eta_\theta} \\
&\quad + T L_\alpha \sqrt{M_\alpha} \frac{\eta_\alpha}{\eta_\theta} + T L_Q \sqrt{M_w} \frac{\eta_w}{\eta_\theta} + T S_\theta M_\theta \eta_\theta \\
&= \frac{\mathbb{E} F(\theta^{(k)}; w_1^{(k)}, \phi^*(\theta^{(k)}), \alpha_0^{(k)}) - \mathbb{E} F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)})}{\eta_\theta} \\
&\quad + \frac{\mathbb{E} F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)}) - \mathbb{E} F(\theta^{(k+1)}; w_{T+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)})}{\eta_\theta} \\
&\quad + T L_\alpha \sqrt{M_\alpha} \frac{\eta_\alpha}{\eta_\theta} + T L_Q \sqrt{M_w} \frac{\eta_w}{\eta_\theta} + T S_\theta M_\theta \eta_\theta, \tag{27}
\end{aligned}$$

where $w_{T+1}^{(k)} = w_1^{(k+1)}$.

For notational simplicity, we define a vector function

$$G(\pi) = \mathbb{E}_{\rho^\pi}[g(\psi_s, \psi_a)].$$

Given a fixed $\theta^{(k)}$, by the definition of G , the optimal

$$\phi^*(\theta^{(k)}) = \frac{1}{\mu}[G(\pi_E) - G(\pi_{\theta^{(k)}})]$$

can be obtained.

Now consider

$$\mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)}) - \mathbb{E}F(\theta^{(k+1)}; w_{T+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)}),$$

by the definition of soft Q-function [17]

$$\begin{aligned} Q_w^{\text{soft}}(s_t, a_t; \theta, \phi) \\ = r_\phi(s_t, a_t) + \mathbb{E}_{(s_{t+1}, \dots) \sim d^{\pi_\theta}} \left[\sum_{l=1}^{\infty} \gamma^l (r_\phi(s_{t+l}, a_{t+l}) + \alpha \mathbb{H}(\pi_\theta(\cdot | s_{t+l}))) \right], \end{aligned}$$

we have

$$\begin{aligned} & \mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)}) - \mathbb{E}F(\theta^{(k+1)}; w_{T+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)}) \\ &= \mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)}) - \mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)}) \\ &= \mathbb{E} \left[\mathbb{E}_{(s,a) \sim \rho^{\pi_E}} [r_{\phi^*(\theta^{(k+1)})}(s, a) - r_{\phi^*(\theta^{(k)})}(s, a)] \right. \\ & \quad + \mathbb{E}_{s_t \sim \mathcal{D}_1^{(k)}} \left[\mathbb{E}_{a_t \sim \pi_{\theta^{(k+1)}}} [-Q_{w_1^{(k+1)}}^{\text{soft}}(s_t, a_t; \theta^{(k+1)}, \phi^*(\theta^{(k+1)})) \right. \\ & \quad \left. \left. + Q_{w_1^{(k+1)}}^{\text{soft}}(s_t, a_t; \theta^{(k+1)}, \phi^*(\theta^{(k)})) \right] \right] - \frac{\mu}{2} (\|\phi^*(\theta^{(k+1)})\|_2^2 - \|\phi^*(\theta^{(k)})\|_2^2) \\ &= \mathbb{E} \left[(\mathbb{E}_{(s,a) \sim \rho^{\pi_E}} [r_{\phi^*(\theta^{(k+1)})}(s, a)] - \mathbb{E}_{(s,a) \sim \rho^{\pi_{\theta^{(k+1)}}}} [r_{\phi^*(\theta^{(k+1)})}(s, a)]) \right. \\ & \quad \left. - (\mathbb{E}_{(s,a) \sim \rho^{\pi_E}} [r_{\phi^*(\theta^{(k)})}(s, a)] - \mathbb{E}_{(s,a) \sim \rho^{\pi_{\theta^{(k+1)}}}} [r_{\phi^*(\theta^{(k)})}(s, a)]) \right] \\ & \quad - \frac{\mu}{2} (\|\phi^*(\theta^{(k+1)})\|_2^2 - \|\phi^*(\theta^{(k)})\|_2^2) \\ &= \mathbb{E} \langle G(\pi_E) - G(\pi_{\theta^{(k+1)}}), \phi^*(\theta^{(k+1)}) - \phi^*(\theta^{(k)}) \rangle \\ & \quad - \frac{\mu}{2} \mathbb{E} \langle \phi^*(\theta^{(k+1)}) + \phi^*(\theta^{(k)}), \phi^*(\theta^{(k+1)}) - \phi^*(\theta^{(k)}) \rangle \\ &= \mathbb{E} \langle \mu \phi^*(\theta^{(k+1)}) - \frac{\mu}{2} (\phi^*(\theta^{(k+1)}) + \phi^*(\theta^{(k)})), \phi^*(\theta^{(k+1)}) - \phi^*(\theta^{(k)}) \rangle \\ &= \frac{\mu}{2} \mathbb{E} \|\phi^*(\theta^{(k+1)}) - \phi^*(\theta^{(k)})\|_2^2 \\ &= \frac{\mu}{2} \mathbb{E} \left\| \frac{1}{\mu} (G(\pi_{\theta^{(k+1)}}) - G(\pi_{\theta^{(k)}})) \right\|_2^2. \end{aligned}$$

Under Assumptions 5 and 4, we get

$$\begin{aligned}
& \mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)}) - \mathbb{E}F(\theta^{(k+1)}; w_{T+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_T^{(k)}) \\
& \leq \frac{S_g^2}{2\mu} \mathbb{E}\|\theta^{(k+1)} - \theta^{(k)}\|_2^2 \leq \frac{TS_g^2}{2\mu} \left(\sum_{t=0}^{T-1} \mathbb{E}\|\theta_{t+1}^{(k)} - \theta_t^{(k)}\|_2^2 \right) \\
& \leq \frac{TS_g^2\eta_\theta^2}{2\mu n_\theta^2} \left(\sum_{t=0}^{T-1} (n_\theta \sum_{j \in D_\theta^t} \mathbb{E}\|\nabla_\theta \tilde{f}^{(j)}(\theta_t^{(k)}; w_{t+1}^{(k)}, \hat{\phi}(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \right) \\
& \leq \frac{T^2 S_g^2 M_\theta \eta_\theta^2}{2\mu}. \tag{28}
\end{aligned}$$

Plugging (28) into (27), we obtain

$$\begin{aligned}
& \sum_{t=0}^{T-1} (\mathbb{E}\|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \\
& \leq \frac{\mathbb{E}F(\theta^{(k)}; w_1^{(k)}, \phi^*(\theta^{(k)}), \alpha_0^{(k)}) - \mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)})}{\eta_\theta} \\
& \quad + TL_\alpha \sqrt{M_\alpha} \frac{\eta_\alpha}{\eta_\theta} + TL_Q \sqrt{M_w} \frac{\eta_w}{\eta_\theta} + (TS_\theta M_\theta + \frac{T^2 S_g^2 M_\theta}{2\mu}) \eta_\theta. \tag{29}
\end{aligned}$$

Summing the equation (29) up, we have

$$\begin{aligned}
& \sum_{k=0}^{N-1} \left(\sum_{t=0}^{T-1} (\mathbb{E}\|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \right) \\
& \leq \frac{1}{\eta_\theta} \sum_{k=0}^{N-1} (\mathbb{E}F(\theta^{(k)}; w_1^{(k)}, \phi^*(\theta^{(k)}), \alpha_0^{(k)}) \\
& \quad - \mathbb{E}F(\theta^{(k+1)}; w_1^{(k+1)}, \phi^*(\theta^{(k+1)}), \alpha_0^{(k+1)})) + NTL_\alpha \sqrt{M_\alpha} \frac{\eta_\alpha}{\eta_\theta} \\
& \quad + NTL_Q \sqrt{M_w} \frac{\eta_w}{\eta_\theta} + N(TS_\theta M_\theta + \frac{T^2 S_g^2 M_\theta}{2\mu}) \eta_\theta.
\end{aligned}$$

Dividing both sides of the above equation by N , we get

$$\begin{aligned}
& \min_{0 \leq k \leq N-1} \sum_{t=0}^{T-1} (\mathbb{E}\|\nabla_\theta F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \\
& \leq \frac{|F(\theta^{(0)}; w_1^{(0)}, \phi^*(\theta^{(0)}), \alpha_0^{(0)}) - \mathbb{E}F(\theta^{(N)}; w_1^{(N)}, \phi^*(\theta^{(N)}), \alpha_0^{(N)})|}{N\eta_\theta} \\
& \quad + TL_\alpha \sqrt{M_\alpha} \frac{\eta_\alpha}{\eta_\theta} + TL_Q \sqrt{M_w} \frac{\eta_w}{\eta_\theta} + (TS_\theta M_\theta + \frac{T^2 S_g^2 M_\theta}{2\mu}) \eta_\theta.
\end{aligned}$$

By Lemma 3, we have

$$|F(\theta^{(0)}; w_1^{(0)}, \phi^*(\theta^{(0)}), \alpha_0^{(0)}) - \mathbb{E}F(\theta^{(N)}; w_1^{(N)}, \phi^*(\theta^{(N)}), \alpha_0^{(N)})| \leq 2B_F.$$

Then we obtain

$$\begin{aligned} I_N &\leq T \min_{0 \leq k \leq N-1} \sum_{t=0}^{T-1} (\mathbb{E} \|\nabla_{\theta} F(\theta_t^{(k)}; w_{t+1}^{(k)}, \phi^*(\theta^{(k)}), \alpha_t^{(k)})\|_2^2) \\ &\leq \frac{2TB_F}{N\eta_{\theta}} + T^2 L_{\alpha} \sqrt{M_{\alpha}} \frac{\eta_{\alpha}}{\eta_{\theta}} + T^2 L_Q \sqrt{M_w} \frac{\eta_w}{\eta_{\theta}} + T^2 M_{\theta} (S_{\theta} + \frac{TS_g^2}{2\mu}) \eta_{\theta}, \end{aligned}$$

where $B_F = \frac{\sqrt{2}\kappa L_g(2-\gamma)}{1-\gamma} + \frac{B_H}{1-\gamma} + \frac{\mu}{2}\kappa^2$.
Given any $\epsilon > 0$, take

$$\begin{aligned} \eta_{\theta} &= \frac{\epsilon}{2T^2 M_{\theta} (2S_{\theta} + TS_g^2/\mu)}, \\ \eta_w &= \frac{\epsilon^2}{8T^4 L_Q M_{\theta} \sqrt{M_w} (2S_{\theta} + TS_g^2/\mu)}, \\ \eta_{\alpha} &= \frac{\epsilon^2}{8T^4 L_{\alpha} M_{\theta} \sqrt{M_{\alpha}} (2S_{\theta} + TS_g^2/\mu)}, \end{aligned}$$

then we need at most

$$N = \frac{8TB_F}{\epsilon\eta_{\theta}} = \tilde{O}\left(\frac{T^3 B_F M_{\theta} (S_{\theta} + TS_g^2/\mu)}{\epsilon^2}\right)$$

iterations such that $I_N \leq \epsilon$. \square

Theorem 3 discloses that despite the min-max computational formulation of GAIL without convex-concave structure, the TSSG algorithm can still converge to a stationary point.

D Experiment Details in Section 6

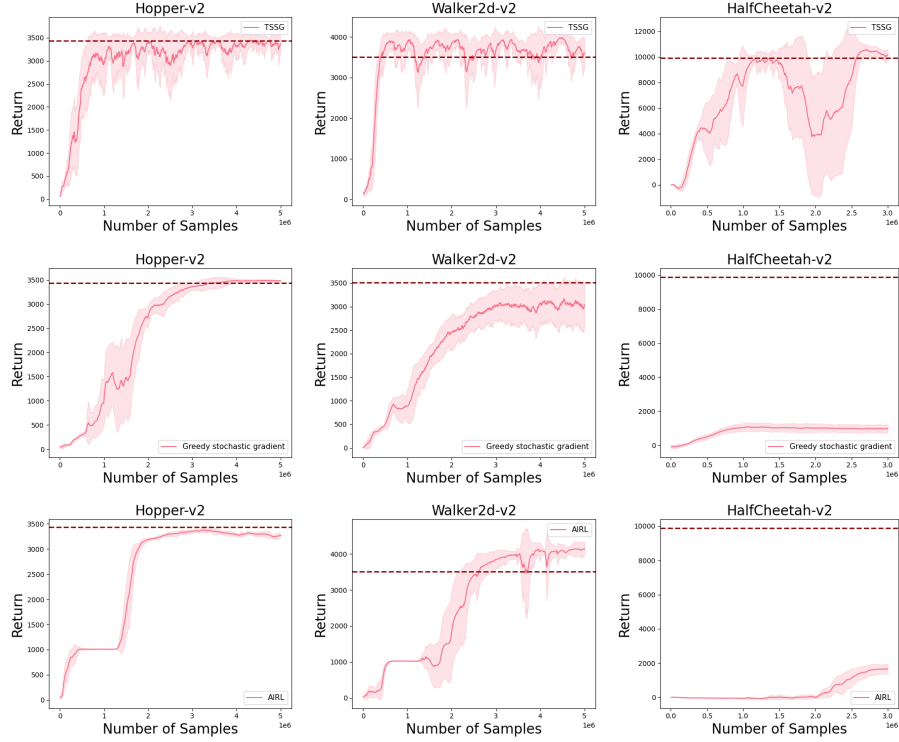
Our experimental environments are MuJoCo, as shown in Table 2, under which three algorithms are evaluated. We use SAC to train the expert policy, and then collect demonstration data with a buffer size of 10^6 . The mean return of the expert demonstrations is given in Table 1. All imitation learning approaches use a two-layer MLP policy network and double two-layer MLP Q-functions. The update frequency and the batch size for TSSG in each environment are listed in Table 3. Every 5000 steps, we evaluate the policy with the mean return of five episodes for all algorithms. Learning curves of the three algorithms are given in Figure 4. The solid lines are the results of the imitation learning approaches and the dashed lines indicate the mean return of expert demonstrations.

Table 2. State dimension, action dimension and max length of the episode in MuJoCo environments.

Environment	State dimension	Action dimension	Max length of the episode
Hopper-v2	11	3	1000
Walker2d-v2	17	6	1000
HalfCheetah-v2	17	6	1000

Table 3. The update frequency and the batch size for TSSG in MuJoCo environments.

Environment	Update frequency	Batch size
Hopper-v2	128	64
Walker2d-v2	128	128
HalfCheetah-v2	128	128

**Fig. 4.** Learning curves comparing TSSG, greedy stochastic gradient and AIRL on Hopper-v2, Walker2d-v2 and HalfCheetah-v2. Each algorithm is run with 5 seeds.